

# Bivariate Regression

Ben Gonzalez

2023-10-06

## Contents

Bivariate Regression Terminology . . . . .	1
Assumptions of Bivariate Regression . . . . .	4
Bivariate Regression Formula . . . . .	4
Scatterplot of the data set . . . . .	4
Calculate the Residual . . . . .	5
Calculate the mean of the Y Values . . . . .	7
Coefficient of Determination or $R^2$ . . . . .	7
Calculate the numerator of the formula - Sum Squared Regression (SSR) . . . . .	7
Calculate the denominator of the formula - Sum Squares Total (SST) . . . . .	8
Calculate the Adjusted-R Squared $Adj.R^2$ or $R^2_{adj}$ . . . . .	8
Utilize the lm function in R to automate our work . . . . .	9
Squaring the correlation $r$ to find the coefficient of determination $R^2$ . . . . .	11
Correlation coefficient of hp and mpg . . . . .	11
Squaring the correlation coefficient . . . . .	12
Data Set Description . . . . .	12


## Bivariate Regression Terminology

- **Regression** – is the process of estimating a best-fitting line that summarizes the relationship between a predictor variable (Independent Variable) and a criterion variable (Dependent Variable).
- **Regression Analysis** – researchers fit a regression line to a sample of data, estimate the parameters of the regression equation (i.e., the constant and regression coefficient), and use the resulting equation to predict scores on a criterion variable.
- **Bivariate** – means that the analyses discussed include just 2 variables, a predictor variable (the X variable), and a criterion variable (the Y variable).
- **Linear** – refers to the fact, when the Y scores are plotted against the X scores, it should be possible to fit a best-fitting straight line through the center of the scores, as opposed to a best-fitting curved line.

## Regression: Generic Analysis Model

---

## PhantomJS not found. You can install it with `webshot::install_phantomjs()`. If it is installed, please



## Assumptions of Bivariate Regression

- **Linearity** – should be able to fit a best-fitting straight line through the scatterplot.
- **Independence** – each observation included in the sample should be drawn independently from the population of interest. Researchers should not have taken repeated measures on the same variable from the same participant.
- **Homogeneity of Variance (Homoscedasticity)** – the variance of the Y scores should remain fairly constant at all values of X.
- **Normality** – residuals of prediction should be normally distributed. Bivariate Normality – for any specific score on one of the variables, scores on the other variable should follow a normal distribution.

## Bivariate Regression Formula

Here we have the formula for the bivariate regression equation. The regression equation takes the following form:

$$\text{Regression Equation : } \hat{y} = a + \beta(X)$$

$\hat{y}$  ~ the predicted score on the criterion variable

$a$  ~ the constant or the intercept of the regression equation.

$\beta$  ~ the unstandardized regression coefficient. Represents the amount of change in Y that is associated with a one-unit

## Scatterplot of the data set

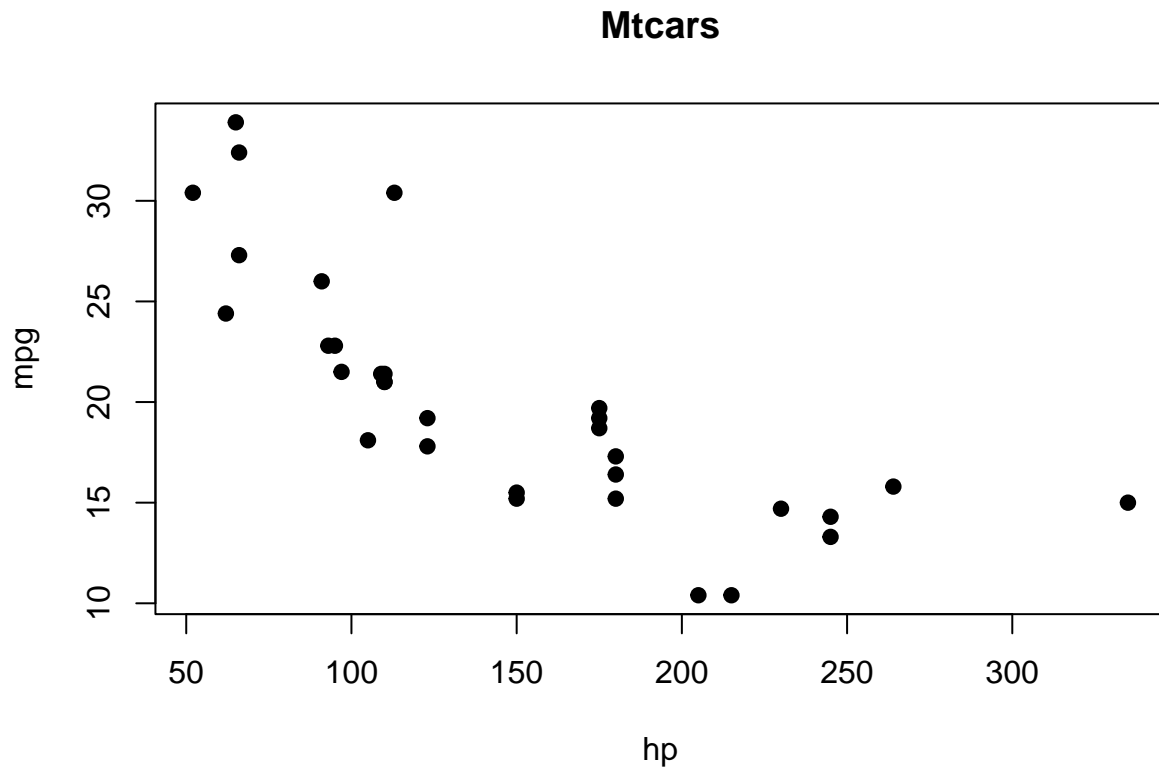
Here we plot our data to get a good look at the shape of the data set.

- **Scatterplot** – a graph that illustrates the nature of the relationship between two quantitative variables.
- **X Axis** – Predictor Variable - hp
- **Y Axis** – Criterion Variable - mpg

We can utilize the following plot function to create a basic scatterplot in R.

```
attach(mtcars)
with(data = mtcars, plot(x = hp, y = mpg, col="black", pch=19, main="Mtcars"))
```

```
## The following object is masked from package:ggplot2:
##
##      mpg
```



### Calculate the Residual

Here we compute the residual by taking the actual  $y$  value and subtract the predicted  $y$  value. The residual for each observation is the difference between the predicted values of  $y$  and the actual values of  $y$ . Calculating the residual helps us to see if we have overpredicted or underpredicted for  $\hat{y}$ .

$$\text{Residual} = \text{actual } y \text{ value} - \text{predicted } y \text{ value}$$

$$r_1 = y_i - \hat{y}_i$$

## [1] "Predicted y Values"

##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
##	22.593750	22.593750	23.753631	22.593750
##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	18.158912	22.934891	13.382932	25.868707
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	23.617174	21.706782	21.706782	17.817770
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	17.817770	17.817770	16.112064	15.429781
##	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
##	14.406357	25.595794	26.550990	25.664022

##	Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
##	23.480718	19.864619	19.864619	13.382932
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
##	18.158912	25.595794	23.890087	22.389065
##	Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
##	12.086595	18.158912	7.242387	22.661978

## [1] "Actual y Values"

## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4  
 ## [16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7  
 ## [31] 15.0 21.4

## [1] "Manually Calculated Residuals"

##	mtcars\$mpg - mpg_prediction\$fitted.values
## Mazda RX4	-1.59374995
## Mazda RX4 Wag	-1.59374995
## Datsun 710	-0.95363068
## Hornet 4 Drive	-1.19374995
## Hornet Sportabout	0.54108812
## Valiant	-4.83489134
## Duster 360	0.91706759
## Merc 240D	-1.46870730
## Merc 230	-0.81717412
## Merc 280	-2.50678234
## Merc 280C	-3.90678234
## Merc 450SE	-1.41777049
## Merc 450SL	-0.51777049
## Merc 450SLC	-2.61777049
## Cadillac Fleetwood	-5.71206353
## Lincoln Continental	-5.02978075
## Chrysler Imperial	0.29364342
## Fiat 128	6.80420581
## Honda Civic	3.84900992
## Toyota Corolla	8.23597754
## Toyota Corona	-1.98071757
## Dodge Challenger	-4.36461883
## AMC Javelin	-4.66461883
## Camaro Z28	-0.08293241
## Pontiac Firebird	1.04108812
## Fiat X1-9	1.70420581
## Porsche 914-2	2.10991276
## Lotus Europa	8.01093488
## Ford Pantera L	3.71340487
## Ferrari Dino	1.54108812
## Maserati Bora	7.75761261
## Volvo 142E	-1.26197823

## [1] "Residual Values"

##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
----	-----------	---------------	------------	----------------

##	-1.59374995	-1.59374995	-0.95363068	-1.19374995
##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	0.54108812	-4.83489134	0.91706759	-1.46870730
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	-0.81717412	-2.50678234	-3.90678234	-1.41777049
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	-0.51777049	-2.61777049	-5.71206353	-5.02978075
##	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
##	0.29364342	6.80420581	3.84900992	8.23597754
##	Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
##	-1.98071757	-4.36461883	-4.66461883	-0.08293241
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
##	1.04108812	1.70420581	2.10991276	8.01093488
##	Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
##	3.71340487	1.54108812	7.75761261	-1.26197823

### Calculate the mean of the Y Values

Here we find the mean of our criterion (y) value of mpg.

$$\bar{y} = \frac{\sum y}{n}$$

We can utilize the **mean** function to calculate the mean of mpg (miles per gallon).

```
mean(mtcars$mpg)
```

```
## [1] 20.09062
```

### Coefficient of Determination or $R^2$

**Coefficient of Determination** – indicates the percent of **variance in the criterion** variable **that is accounted for by the predictor** variable.

---


$$\text{Coefficient of Determination : } R^2 = 1 - \frac{\text{sum squared regression (SSR)}}{\text{sum squares total (SST)}}$$

$$= 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2} \quad y_i = \text{actual } y \text{ values} \quad \hat{y}_i = \text{predicted } y \text{ values} \quad \bar{y} = \text{mean of } y \quad \sum \text{ or sigma} = \text{sum}$$


---

### Calculate the numerator of the formula - Sum Squared Regression (SSR)

$$\sum (y_i - \hat{y}_i)^2$$

```
top_of_formula <- sum(mpg_prediction$residuals^2)
print(top_of_formula)
```

```
## [1] 447.6743
```

Calculate the denominator of the formula - Sum Squares Total (SST)

---


$$\sum (y_i - \bar{y})^2$$


---

```
bottom_of_formula <- sum((mtcars$mpg - mean(mtcars$mpg))^2)
print(bottom_of_formula)
```

```
## [1] 1126.047
```

---


$$R^2 = 1 - \frac{447.6743}{1126.047}$$

$$R^2 = 1 - 0.3975627$$

$$R^2 = 0.6024373 \quad R^2 = 0.6024$$


---

Calculate the Adjusted-R Squared  $Adj.R^2$  or  $R^2_{adj}$

---


$$Adj.R^2 \text{ or } R^2_{adj} = 1 - (1 - R^2) \cdot (n-1)/(n-p-1) \quad Adj.R^2 \text{ or } R^2_{adj} = 1 - (1 - 0.6024373) \cdot (32-1)/(32-1-1) R^2 = \text{coefficient of } a$$


---

```
adj.r.squared = 1 - (1 - 0.6024373) * ((32 - 1)/(32-1-1))
print(adj.r.squared)
```

```
## [1] 0.5891852
```



## Utilize the `lm` function in R to automate our work

Here we can utilize the ***lm*** function in R to perform our bivariate regression (simple linear regression). This will allow us to save the model to a variable and then utilize the **\$** (dollar sign) operator in R. The **\$** (dollar sign) operator allows us to pull out things we need such as the residuals and fitted values that are returned from the summary function.

```
mpg_hp_model <- lm(mpg ~ hp, mtcars)
print(mpg_hp_model$residuals)
```

### Print the residuals of the model

##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
##	-1.59374995	-1.59374995	-0.95363068	-1.19374995
##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	0.54108812	-4.83489134	0.91706759	-1.46870730
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	-0.81717412	-2.50678234	-3.90678234	-1.41777049
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	-0.51777049	-2.61777049	-5.71206353	-5.02978075
##	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
##	0.29364342	6.80420581	3.84900992	8.23597754
##	Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
##	-1.98071757	-4.36461883	-4.66461883	-0.08293241
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
##	1.04108812	1.70420581	2.10991276	8.01093488
##	Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
##	3.71340487	1.54108812	7.75761261	-1.26197823

```
mpg_hp_model <- lm(mpg ~ hp, mtcars)
print(mpg_hp_model$coefficients)
```

### Print the coefficients of the model

```
## (Intercept)      hp
## 30.09886054 -0.06822828
```

```
mpg_hp_model <- lm(mpg ~ hp, mtcars)
print(mpg_hp_model$fitted.values)
```

### Print the fitted values of the model

##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
##	22.593750	22.593750	23.753631	22.593750

##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	18.158912	22.934891	13.382932	25.868707
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	23.617174	21.706782	21.706782	17.817770
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	17.817770	17.817770	16.112064	15.429781
##	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
##	14.406357	25.595794	26.550990	25.664022
##	Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
##	23.480718	19.864619	19.864619	13.382932
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
##	18.158912	25.595794	23.890087	22.389065
##	Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
##	12.086595	18.158912	7.242387	22.661978

**Putting it altogether** Here we can print out the summary of the model utilizing the *summary* function in R; `summary(mpg_hp_model)`. We can also plot the predicted y values with the actual y values. Then we can draw a line between each of the predicted values and the actual values. This helps us visualize the amount of variation that is present between the predicted vs the actual values of y.


```
mpg_hp_model = lm(mpg ~ hp, mtcars)
print(summary(mpg_hp_model))
```

**Plot our residuals and a best fitting line** Here we can utilize the *ggplot2* package to plot our model. We can also plot the residuals along with a best fitting line.

```
mtcars %>% ggplot(aes(hp,mpg))+
  geom_point()+
  geom_smooth(method = "lm")+
  geom_linerange(aes(ymax = mpg, ymin = mpg-resid),color="red")
```

```
##
## Call:
## lm(formula = mpg ~ hp, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.7121 -2.1122 -0.8854  1.5819  8.2360
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  30.09886    1.63392   18.421  < 2e-16 ***
## hp          -0.06823    0.01012   -6.742  1.79e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.863 on 30 degrees of freedom
## Multiple R-squared:  0.6024, Adjusted R-squared:  0.5892
## F-statistic: 45.46 on 1 and 30 DF, p-value: 1.788e-07

## 'geom_smooth()' using formula = 'y ~ x'
```



Now we can take a predictor value ( $X$ ) and plug it in. We then are able to predict where our criterion value ( $Y$ ) will be.

$$\hat{y} = 30.09886 + -0.06823(X)$$

$$\hat{y} = 30.09886 + -0.06823(335)$$

$$\hat{y} = 30.09886 + -22.85705$$

$$\hat{y} = 7.24$$

### Squaring the correlation $r$ to find the coefficient of determination $R^2$

According to Hatcher (2013) we can simply square the correlation provided we are looking at only one predictor variable and one dependent variable. When we square the correlation coefficient this will give us the coefficient of determination.

### Correlation coefficient of hp and mpg

```
cor(x = mtcars$hp,mtcars$mpg)
```

```
## [1] -0.7761684
```

$r = -0.7761684$

### Squaring the correlation coefficient

Here we can square the correlation coefficient  $r$  and it will give us the coefficient of determination or  $R^2$

```
cor(x = mtcars$hp,mtcars$mpg)^2
```

```
## [1] 0.6024373
```

$R^2 = 0.6024373$

Here we can see we get the same value for the coefficient of determination  $R^2$  by squaring the correlation as if we had utilized the **lm** function. However the **lm** function has advantages as it provides us with our p-value, F-statistic, and the intercept and the unstandardized regression coefficient.

### Data Set Description

Motor Trend Car Road Tests

Description:

The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models).

References

Hatcher, L. (2013). *Advanced statistics in research: Reading, understanding, and writing up data analysis results*. Shadow Finch Media.

Henderson and Velleman (1981). dataset: Motor Trend Car Road Tests. R package version 4.3.1

Henderson and Velleman (1981), Building multiple regression models interactively. Biometrics, 37, 391–411.